

## Assignment – 2

### SMS Spam Classification

Assignment Date	08 October 2022
Student Name	Rohit Jacob George
Student Roll Number	2019504572
Maximum Marks	2 Marks

## Task - 1: Import the necessary libraries

### 1.Import the Necessary Libraries

```
In [2]: import numpy as np
import pandas as pd
import os
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

/kaggle/input/sms-spam-collection-dataset/spam.csv

## Task - 2: Reading the .csv dataset

### 2. Reading the .csv dataset

```
In [3]: data=pd.read_csv("../input/sms-spam-collection-dataset/spam.csv",encoding="latin")
data.head()
```

```
Out[3]:
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

```
In [4]: data.columns
```

```
Out[4]: Index(['v1', 'v2', 'Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], dtype='object')
```

## Task - 3: Drop the unnamed columns

### 3. Drop the unnamed Columns

```
In [5]: data=data.drop(columns=["Unnamed: 2","Unnamed: 3","Unnamed: 4"])
```

## Task – 4: Renaming Column names sensible

### 4. Renaming Column names sensible

```
In [6]: data=data.rename(  
        {  
            "v1":"Category",  
            "v2":"Message"  
        },  
        axis=1  
    )
```

```
In [7]: data.head()
```

```
Out[7]:
```

	Category	Message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

## Task – 5: Check for null values in dataset

### 5. Check for null values in dataset

```
In [8]: data.isnull().sum()
```

```
Out[8]: Category    0  
        Message     0  
        dtype: int64
```

```
In [9]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5572 entries, 0 to 5571  
Data columns (total 2 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   Category    5572 non-null   object  
1   Message     5572 non-null   object  
dtypes: object(2)  
memory usage: 87.2+ KB
```


## Task – 6: Creating a new field to store the Message Lengths

### 6.Creating a new Field to store the Message Lengths

```
In [10]: data["Message Length"]=data["Message"].apply(len)
```

## Task – 7: Histogram Inference of Message Lengths of Spam and Non-spam messages

```
In [11]: fig=plt.figure(figsize=(12,8))
sns.histplot(
    x=data["Message Length"],
    hue=data["Category"]
)
plt.title("ham & spam messege length comparision")
plt.show()
```



```
In [12]: ham_desc=data[data["Category"]=="ham"]["Message Length"].describe()
spam_desc=data[data["Category"]=="spam"]["Message Length"].describe()

print("Ham Messege Length Description:\n",ham_desc)
print("*****")
print("Spam Message Length Description:\n",spam_desc)
```

```
Ham Messege Length Description:
count    4825.000000
mean      71.023627
std       58.016023
min        2.000000
25%       33.000000
50%       52.000000
75%       92.000000
max      910.000000
Name: Message Length, dtype: float64
*****
Spam Message Length Description:
count      747.000000
mean     138.866131
std      29.183082
min      13.000000
25%     132.500000
50%     149.000000
75%     157.000000
max     224.000000
Name: Message Length, dtype: float64
```

```
In [13]: data.describe(include="all")
```

```
Out[13]:
```

	Category	Message	Message Length
count	5572	5572	5572.000000
unique	2	5169	NaN
top	ham	Sorry, I'll call later	NaN
freq	4825	30	NaN
mean	NaN	NaN	80.118808
std	NaN	NaN	59.690841
min	NaN	NaN	2.000000
25%	NaN	NaN	36.000000
50%	NaN	NaN	61.000000
75%	NaN	NaN	121.000000
max	NaN	NaN	910.000000


## Task – 8: Visualizing count of messages of Spam and Non Spam

### 8. Visualizing count of messages of Spam and Non Spam

```
In [14]: data["Category"].value_counts()
```

```
Out[14]: ham      4825  
spam       747  
Name: Category, dtype: int64
```

```
In [15]: sns.countplot(  
    data=data,  
    x="Category"  
)  
plt.title("ham vs spam")  
plt.show()
```



```
In [16]: ham_count=data["Category"].value_counts()[0]  
spam_count=data["Category"].value_counts()[1]  
  
total_count=data.shape[0]  
  
print("Ham contains:{:.2f}% of total data.".format(ham_count/total_count*100))  
print("Spam contains:{:.2f}% of total data.".format(spam_count/total_count*100))
```

```
Ham contains:86.59% of total data.  
Spam contains:13.41% of total data.
```

## Task – 9: Undersampling to Genralize Model and Balance Spam and Ham quantities in dataset

## 9. Undersampling to Genralize Model and Balance Spam and Ham quantities in dataset

```
In [17]: minority_len=len(data[data["Category"]=="spam"])
majority_len=len(data[data["Category"]=="ham"])
minority_indices=data[data["Category"]=="spam"].index
majority_indices=data[data["Category"]=="ham"].index
random_majority_indices=np.random.choice(
    majority_indices,
    size=minority_len,
    replace=False
)

undersampled_indices=np.concatenate([minority_indices,random_majority_indices])
df=data.loc[undersampled_indices]
df=df.sample(frac=1)
df=df.reset_index()
df=df.drop(
    columns=["index"],
)
```


```
In [18]: df.shape
```

```
Out[18]: (1494, 3)
```

```
In [19]: df["Category"].value_counts()
```

```
Out[19]: ham      747
spam      747
Name: Category, dtype: int64
```

```
In [20]: sns.countplot(
    data=df,
    x="Category"
)
plt.title("ham vs spam")
plt.show()
```



Display the head of new **df**

```
In [21]: df.head()
```

```
Out[21]:
```

	Category	Message	Message Length
0	spam	FREE>Ringtone! Reply REAL or POLY eg REAL1 1. ...	158
1	spam	URGENT! We are trying to contact U Todays draw...	157
2	ham	Ok ill send you with in &lt;DECIMAL> ok.	45
3	ham	Oh just getting even with u.... u?	34
4	spam	A link to your picture has been sent. You can ...	96

## Task – 10: Binary Encoding of Spam and Ham Categories

```
In [22]: df["Label"]=df["Category"].map(
        {
            "ham":0,
            "spam":1
        }
    )
```

```
In [23]: df.head()
```

```
Out[23]:
```

	Category	Message	Message Length	Label
0	spam	FREE>Ringtone! Reply REAL or POLY eg REAL1 1. ...	158	1
1	spam	URGENT! We are trying to contact U Todays draw...	157	1
2	ham	Ok ill send you with in &lt;DECIMAL>&gt; ok.	45	0
3	ham	Oh just getting even with u.... u?	34	0
4	spam	A link to your picture has been sent. You can ...	96	1

## Task – 11: Import Necessary Libraries to perform Word Tokenization

```
In [24]: import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

stemmer=PorterStemmer()
```

```
In [25]: corpus=[]
for message in df["Message"]:
    message=re.sub("[^a-zA-Z]", " ",message)
    message=message.lower()
    message=message.split()
    message=[stemmer.stem(words)
              for words in message
              if words not in set(stopwords.words("english"))
            ]
    message=" ".join(message)
    corpus.append(message)
```

## Task – 12: Perform One Hot on Corpus


```
In [26]: from tensorflow.keras.preprocessing.text import one_hot
vocab_size=10000

oneHot_doc=[one_hot(words,n=vocab_size)
             for words in corpus
            ]
```

```
In [27]: df["Message Length"].describe()
```

```
Out[27]: count    1494.000000
mean         104.491299
std           60.362332
min            2.000000
25%           49.000000
50%          114.000000
75%          153.000000
max           910.000000
Name: Message Length, dtype: float64
```

```
In [28]: fig=plt.figure(figsize=(12,8))
sns.kdeplot(
    x=df["Message Length"],
    hue=df["Category"]
)
plt.title("ham & spam messege length comparision")
plt.show()
```



```
In [29]: from tensorflow.keras.preprocessing.sequence import pad_sequences
sentence_len=200
embedded_doc=pad_sequences(
    oneHot_doc,
    maxlen=sentence_len,
    padding="pre"
)
```

```
In [30]: extract_features=pd.DataFrame(
    data=embedded_doc
)
target=df["Label"]
```

```
In [31]: df_final=pd.concat([extract_features,target],axis=1)
```

```
In [32]: df_final.head()
```

```
Out[32]:
```

	0	1	2	3	4	5	6	7	8	9	...	191	192	193	194	195	196	197	198	199	Label
0	0	0	0	0	0	0	0	0	0	0	...	8116	8983	7883	1884	5957	5877	266	1527	5846	1
1	0	0	0	0	0	0	0	0	0	0	...	9989	7682	5710	5519	2447	1240	3994	6950	3655	1
2	0	0	0	0	0	0	0	0	0	0	...	0	0	3310	6099	7761	9276	4679	2205	3310	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	8194	7945	3841	266	266	0
4	0	0	0	0	0	0	0	0	0	0	...	5677	7440	8481	9975	2366	4841	4320	4320	4672	1

5 rows × 201 columns

## Task – 13: Splitting Dependent and Independent Variables

```
In [33]: X=df_final.drop("Label",axis=1)
y=df_final["Label"]
```

## Task – 14: Train, test and Validation Split

```
In [34]: from sklearn.model_selection import train_test_split
```

```
In [35]: X_trainval,X_test,y_trainval,y_test=train_test_split(
        X,
        y,
        random_state=42,
        test_size=0.15
    )
```

```
In [36]: X_train,X_val,y_train,y_val=train_test_split(
        X_trainval,
        y_trainval,
        random_state=42,
        test_size=0.15
    )
```

## Task – 15: Building a Sequential Model

```
In [37]: from tensorflow.keras.layers import LSTM
        from tensorflow.keras.layers import Dense
        from tensorflow.keras.layers import Embedding
        from tensorflow.keras.models import Sequential
```

```
In [38]: model=Sequential()
```

```
In [39]: feature_num=100
        model.add(
            Embedding(
                input_dim=vocab_size,
                output_dim=feature_num,
                input_length=sentence_len
            )
        )
        model.add(
            LSTM(
                units=128
            )
        )
        model.add(
            Dense(
                units=1,
                activation="sigmoid"
            )
        )
        model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
-----		
embedding (Embedding)	(None, 200, 100)	1000000
-----		
lstm (LSTM)	(None, 128)	117248
-----		
dense (Dense)	(None, 1)	129
-----		
Total params: 1,117,377		
Trainable params: 1,117,377		
Non-trainable params: 0		
-----		



```
In [40]: from tensorflow.keras.optimizers import Adam
model.compile(
    optimizer=Adam(
        learning_rate=0.001
    ),
    loss="binary_crossentropy",
    metrics=["accuracy"]
)
```


## Task – 16: Model Fitting

```
In [41]: history=model.fit(
    X_train,
    y_train,
    validation_data=(
        X_val,
        y_val
    ),
    epochs=10
)
```

```
Epoch 1/10
34/34 [=====] - 24s 633ms/step - loss: 0.6324 - accuracy: 0.6331 - val_loss: 0.4218 - val_accuracy: 0.8377
Epoch 2/10
34/34 [=====] - 21s 608ms/step - loss: 0.3045 - accuracy: 0.9257 - val_loss: 0.1631 - val_accuracy: 0.9529
Epoch 3/10
34/34 [=====] - 21s 609ms/step - loss: 0.1046 - accuracy: 0.9689 - val_loss: 0.1231 - val_accuracy: 0.9581
Epoch 4/10
34/34 [=====] - 21s 621ms/step - loss: 0.0465 - accuracy: 0.9880 - val_loss: 0.1293 - val_accuracy: 0.9581
Epoch 5/10
34/34 [=====] - 21s 613ms/step - loss: 0.0342 - accuracy: 0.9895 - val_loss: 0.1252 - val_accuracy: 0.9686
Epoch 6/10
34/34 [=====] - 21s 615ms/step - loss: 0.0179 - accuracy: 0.9951 - val_loss: 0.1366 - val_accuracy: 0.9686
Epoch 7/10
34/34 [=====] - 21s 614ms/step - loss: 0.0121 - accuracy: 0.9968 - val_loss: 0.1314 - val_accuracy: 0.9634
Epoch 8/10
34/34 [=====] - 21s 619ms/step - loss: 0.0222 - accuracy: 0.9944 - val_loss: 0.1479 - val_accuracy: 0.9634
Epoch 9/10
34/34 [=====] - 21s 614ms/step - loss: 0.0077 - accuracy: 0.9989 - val_loss: 0.1624 - val_accuracy: 0.9634
Epoch 10/10
34/34 [=====] - 21s 614ms/step - loss: 0.0077 - accuracy: 0.9976 - val_loss: 0.1751 - val_accuracy: 0.9634
```

```
In [42]: metrics = pd.DataFrame(history.history)
metrics.rename(columns = {'loss': 'Training_Loss', 'accuracy': 'Training_Accuracy', 'val_loss': 'Validation_Loss', 'val_accuracy':
def plot_graph_acc(var1, var2, string):
    metrics[[var1, var2]].plot()
    plt.title('Training and Validation ' + string)
    plt.xlabel('Number of epochs')
    plt.ylabel(string)
    plt.legend([var1, var2])
```

```
In [43]: plot_graph_acc('Training_Accuracy', 'Validation_Accuracy', 'accuracy')
```



```
In [44]: y_pred=model.predict(X_test)
y_pred=(y_pred>0.5)
```

```
In [45]: model.save('Spam_SMS_classifier.h5')
```

## Task – 16: Evaluating the Model


```
In [46]: from sklearn.metrics import accuracy_score, confusion_matrix
```

```
In [47]: score=accuracy_score(y_test,y_pred)
print("Test Score:{:.2f}%".format(score*100))
```

Test Score:96.89%

```
In [48]: cm=confusion_matrix(y_test,y_pred)
fig=plt.figure(figsize=(12,8))
sns.heatmap(
    cm,
    annot=True,
)
plt.title("Confusion Matrix")
cm
```


```
Out[48]: array([[104,  0],
               [ 7, 114]])
```



## Task – 18: Function to Test the Model on a Random message

```
In [62]: def classify_message(model,message):
    for sentences in message:
        sentences=nlk.sent_tokenize(message)
        for sentence in sentences:
            words=re.sub("[^a-zA-Z]", " ", sentence)
            if words not in set(stopwords.words('english')):
                word=nlk.word_tokenize(words)
                word=" ".join(word)
    oneHot=[one_hot(word,n=vocab_size)]
    text=pad_sequences(oneHot,maxlen=sentence_len,padding="pre")
    predict=model.predict(text)
    if predict>0.5:
        print("It is a spam")
        print("predict score: ", predict[0][0])
    else:
        print("It is not a spam")
        print("predict score: ", predict[0][0])
```

```
In [80]: message1="I am having my Tests right now. Will call back as soon as possible! Till then be safe wherever you are. Be Alert of any  
message2="Your Rs.8850 welcome bonus is ready to be credited. Download Jungle Rummy now. Claim Bonus on your first deposit prize
```



```
In [81]: classify_message(model,message1)
```

```
It is not a spam  
predict score: 0.037389785
```

```
In [82]: classify_message(model,message2)
```

```
It is a spam  
predict score: 0.9936712
```